

Networking

How to use it in Claude, ChatGPT, and Codex

Three skills that research a company, map the people inside it, and turn public activity into engagement tips. They read what is public, never post or message, and never write your outreach.

1. Pick an intent

Everything these skills produce is shaped by one input: your intent. Five presets ship with them, and anything else you describe in your own words works too, in which case the skill builds a matching profile, states it back, and proceeds once you confirm.

Intent	The question it answers
job-search	Is this a good moment to pursue this team, and who decides?
client-prospecting	Is this a live opportunity, what triggered it, and who owns the budget?
partner-bd	What partnership fits how they already partner, and is the slot taken?
investor-outreach	Does the thesis fit, are they deploying, and is there a conflict?
talent-sourcing	Is this team stable or loosening, and where is the movable talent?

Your last intent is remembered and reused as the default. The skill says when it is falling back to one, so a stale default is easy to catch.

2. Start a run

Host	Company research	Contact research
Claude	/networking:research-companies	/networking:research-contacts
ChatGPT	@networking-research-companies	@networking-research-contacts
Codex	\$networking-research-companies	\$networking-research-contacts

Plain language works everywhere and is usually faster:

```
Research Acme Corp for client prospecting, engineering team, Dubai office
```

You get a chat answer first, then a Word document named for the subject and timestamped in your own zone. Judge the research in the chat; keep the document.

3. Use it inside a project

A project is where these skills become yours rather than generic, because they read the project's files as context on every run. Create one, add the skills, and put two files in it:

- `intent-defaults.md`: your standing context and default intent
- `learned-preferences.md`: the rules you have taught the skills

Anything else you drop in gets used too, and this is the part people miss. Add your CV and briefs are judged against your actual background. Add a page on what you sell and prospecting runs can say whether a company is a real fit rather than listing signals and leaving the judgement to you. Add accounts to avoid and they stay out of shortlists.

Say so once at the start of a run, and the skill will draw on them:

```
Use my CV and the ICP note in this project when you research Acme
```

4. The feedback loop

When an output misses, say so in the words you would use to a colleague, and the correction becomes permanent.

```
/networking:feedback the shortlist was full of VPs who can't actually buy anything
```

In ChatGPT say `@networking-feedback` and then the same sentence. What happens next never varies: it separates a lasting preference from a one-off factual error, drafts only the first as a specific testable rule, gives it a scope, shows you the exact text before saving, and says what changes next run.

Every rule carries one scope, and the narrowest that fits is the right one:

- `global`: every run. Use sparingly.
- `intent:client-prospecting`: only under that intent. Most rules belong here.
- `skill:research-contacts`: only to that skill, whatever the intent.

Narrow scopes age better: a complaint about one prospecting shortlist is rarely a rule about everything you do. Conflicting rules are superseded rather than deleted, so the file still explains why the skills behave as they do.

Be specific. “The tips were bad” gives it nothing; “stop guessing email addresses, leave the cell blank” becomes a rule you never repeat.

Where the rule is saved

In Claude and Codex it is written to your memory folder automatically. In ChatGPT, where a skill has no writable storage, it shows the block and asks you to paste it into `learned-preferences.md` in your project files. It never claims a rule was saved when it was only displayed, so if you skip the paste, the rule is gone.

5. Fill the memory file

Rules tell the skills how to behave. `intent-defaults.md` tells them who you are, which is often worth more.

```
last_intent: client-prospecting
timezone: Asia/Dubai
```

Below the settings, fill in what you do and what you sell, the roles or functions you target, the markets you cover, and any companies to always exclude.

6. Import what your AI already remembers

If you have worked with ChatGPT or Claude a while, it already knows things worth carrying over:

```
Print everything you remember about my work, my clients, and how I like outputs
```

Paste the result into the feedback skill and ask it to sort the pile. Durable facts go into `intent-defaults.md` as standing context; instructions about output become scoped rules. You see both blocks first, so you can drop whatever is stale.

```
Turn this into standing context and rules for the networking skills
```

7. Where the research actually comes from

There are no built-in data feeds, no contact database, and no paid API behind any of this. The skills read the open web through whatever your assistant can already reach, which means the quality of a brief depends on which of three capabilities is available to it.

Capability	What it does	Without it
Web search	Finds news, funding coverage, press releases, company pages, filings	The skills cannot run. This is the floor.
Page fetching	Reads the pages search returns, in full rather than as a snippet	Claims get thinner and dates go missing
Browsing	Opens JavaScript-heavy pages: LinkedIn profiles, People tabs, social timelines	People mapping and activity research degrade sharply; the brief says so

Browsing is the one that changes results most. In Claude that is the Claude in Chrome extension; in ChatGPT it is the built-in browsing or agent mode. Without it the skills fall back to search results alone and flag every gap rather than filling it with guesswork.

Company research: what it looks at

Source	What is pulled from it
News coverage	Hiring and layoffs, acquisitions, leadership changes, expansion, restructuring
Press releases and company newsroom	Announcements in the company's own words, with reliable dates
Company website: about, careers, team pages	Which departments exist in which locations, open-role volume
Company blog and engineering blog	What teams are actually building, technical direction
Public filings	Ownership, financial position, disclosed risks, named officers
Funding coverage surfaced through search	Round, amount, lead investors, and which function the money was tied to
LinkedIn company page and People tab	Team structure, headcount direction, named people and titles
Earnings call coverage and transcripts, where public	Stated priorities and publicly admitted problems, useful for prospecting
Reviews and public commentary (Glassdoor-style)	Culture and morale signals, treated as sentiment and never as fact
Procurement notices and tenders, where public	Live buying activity, used only under commercial intents
Partnership, integration, and marketplace listings	How the company partners, and which slots are already filled

Which of these get weight depends on your intent. A job-search run leans on careers pages, hiring news, and team growth. A client-prospecting run leans on funding, fiscal timing, new executives, vendor switches, and publicly stated pain. An investor-outreach run leans on recent deals and portfolio adjacency. The source list is the same; the emphasis is not.

Contact research: what it looks at

Source	What is pulled from it
LinkedIn profile and activity	Posts, articles, comments, reactions, role changes, recurring themes
X (x.com)	Posts and replies, topics, tone, how much they engage versus broadcast
Threads	The same, where the person is active there
Conference and speaker pages	Talks given, topics they are invited to speak on, sometimes a work email
Podcast appearances	Longer-form views that rarely fit in a post
Bylines and newsletters	What they write at length, and how they argue
Public code repositories	What they build, how active they are, occasionally a contact address
Company website and team pages	Confirmed title, location, and sometimes a published work email
Post timestamps across all of the above	The estimate of which days and hours they are most active

Work emails come only from places a company or the person published them: a website, a byline, a speaker page, a repository. An address worked out from a pattern is labelled *inferred*, *unconfirmed* and never presented as *verified*. If nothing public exists, the cell reads *not found*, which is the honest answer.

How much each source actually yields

Two of the most useful sources are also the most restricted, and it is better to know that before you read a thin section and assume the research was lazy.

Source	Logged-out access	What to expect
LinkedIn	Heavily gated	Profiles are often partly visible; the People tab usually needs a session. Browsing helps; some profiles still cannot be read.
X (x.com)	Restricted	Individual posts surfaced by search are usually reachable, full timelines frequently are not. Coverage is uneven and the brief says when it is thin.
Threads	Mostly open	Public posts generally readable, though the account has to actually be active there.
News, press, filings	Open	The most reliable layer, and where dated claims come from.
Paywalled publications	Blocked	Not bypassed. If a claim only exists behind a paywall, it is left out or marked <i>unconfirmed</i> .

Since X coverage is the least predictable, treat a sparse X section as a fact about access rather than about the person. Where activity data is too thin to support an estimate, the skills say so instead of inventing a posting window.

What it never touches

- Contact databases and data brokers. No ZoomInfo, Apollo, Clearbit, Sales Navigator, or equivalent, and no purchased lists.
- Email-finder and verification APIs. Addresses come from published pages or they do not appear.
- Anything behind a login you do not legitimately hold, and anything behind a paywall.
- Private or personal material that happens to be technically findable: home address, family, health, personal accounts unrelated to work.
- Your own inbox, CRM, or calendar, unless you deliberately connect one.
- Any action on a platform. No connection requests, messages, follows, reactions, or form submissions.

How to tell what was used

Momentum research prioritises the last twelve months and dates every claim, so you can judge staleness yourself. Anything unverified is labelled in the same cell rather than in a disclaimer at the bottom. And every brief ends with a confidence note naming what could not be confirmed, which profiles could not be read, and whether browsing was available for that run, which is what lets you tell a thin result from a thin company.

8. Giving it more to work with

Two ways to widen what the research can reach, both under your control.

- Project files. Anything you add becomes readable context: your CV, an ICP description, a target account list, past call notes, a competitor teardown. This is the highest-leverage change, because it lets the skills judge fit rather than just report signals.
- Connectors. If you connect a search tool, a CRM, or a data source through your assistant, the skills will use it alongside the open web. Nothing is connected by default, so this is always something you switched on.

Two rules that never bend

You get tips, never the message: what someone posts about, how they engage, when they are active, what angle is natural. You write the outreach, because research scales and a message that sounds like you does not. This is switchable for a single intent through the feedback loop, never globally and never by asking mid-conversation, so turning it on is always deliberate.

And nothing is invented. Anything unverified is labelled rather than stated as fact, people and titles are never made up to fill a table, and if someone has publicly said they do not want to be approached, the brief says so and advises against it whatever your intent.